Open in app
Sign up
Medium Logo
Search
Write

Cloud AI Platforms: AWS vs. GCP vs. Azure for Machine-Learning Workloads

Vitaly Sem

5 min read

Nov 6, 2025

Audience: engineering managers, machine-learning (ML) leads, founders validating an ML stack.

Note: Pricing and limits here are illustrative. They change often — always check the current docs and your region.

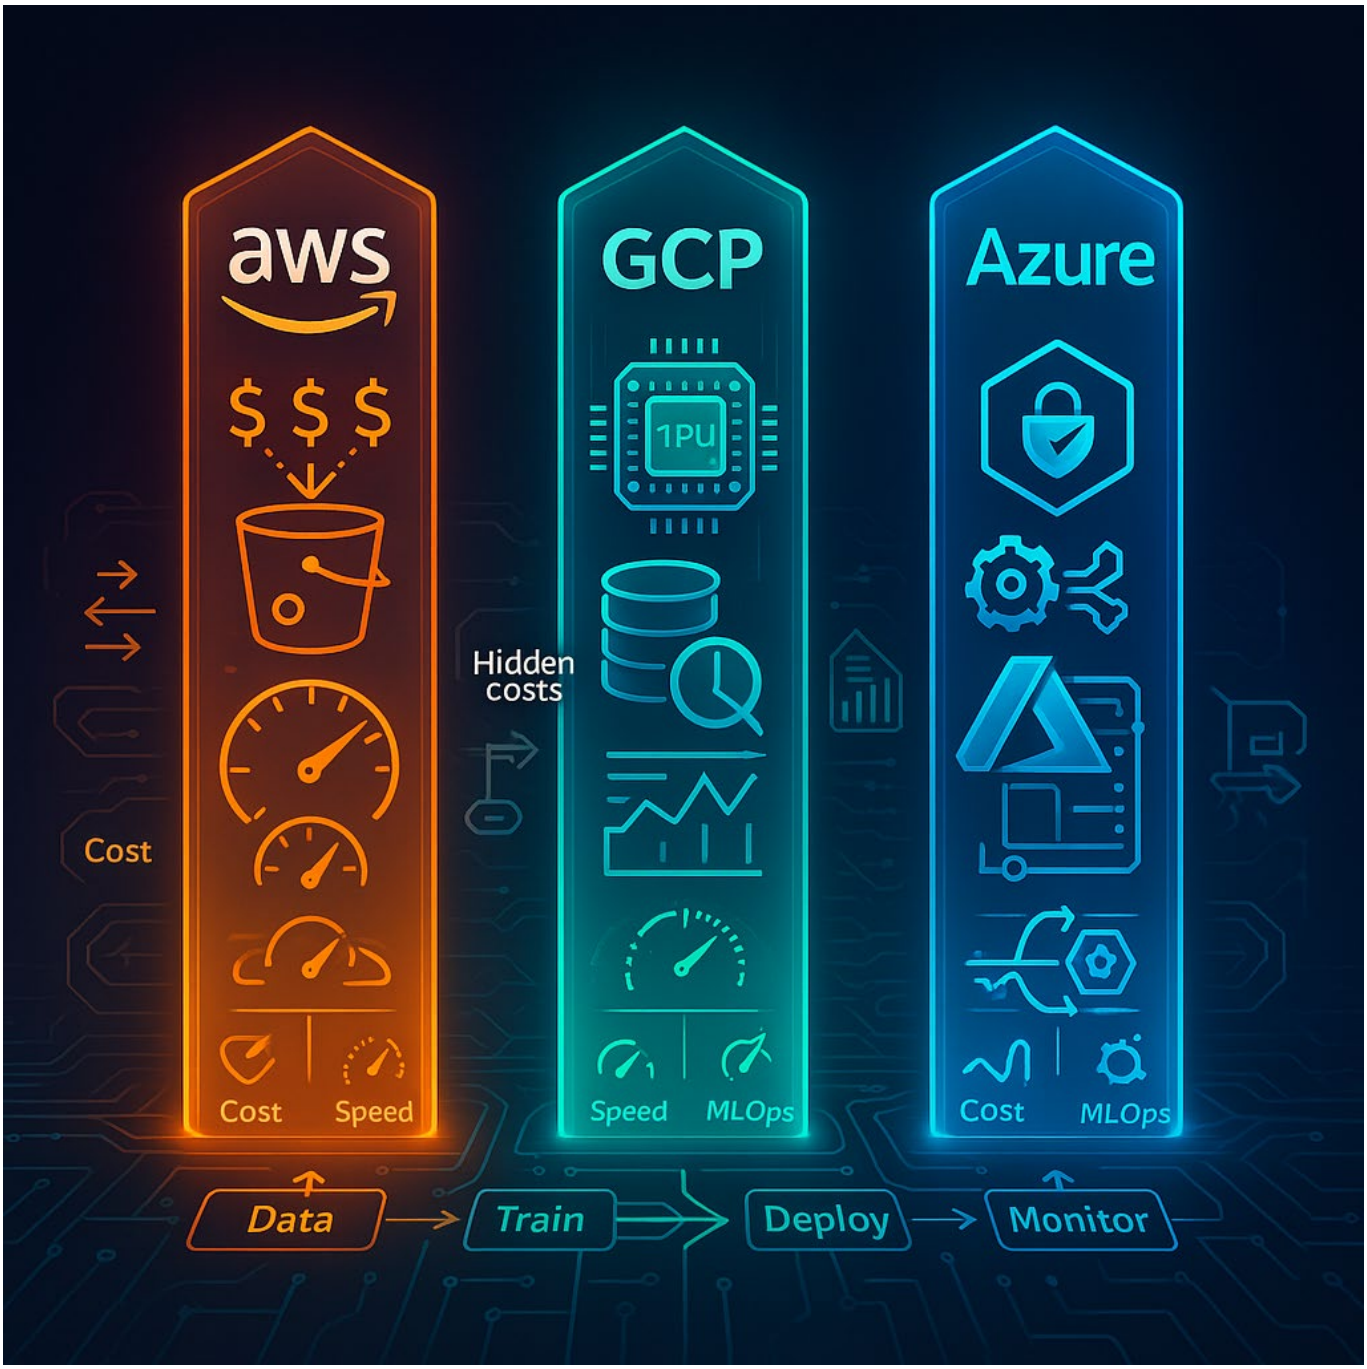

How we evaluate

We compare the platforms across three practical dimensions:

- **Latency & throughput** on like-for-like GPU (Graphics Processing Unit) setups
- **End-to-end cost** of training and serving across different data scales
- **MLOps (Machine Learning Operations) integration**: deployment,

monitoring, and model lifecycle tooling

Quick prototyping tip: **BotHub** is an aggregator of AI models that hides API-key wrangling and setup, so you can focus on experiments instead of plumbing. New accounts get promo credits for test runs.

Amazon SageMaker (AWS — Amazon Web Services)

A fully managed service covering the ML lifecycle — from data prep and experimentation to deployment, monitoring, and governance.

Pricing posture

AWS markets a **pay-for-what-you-use** model across compute, storage, and data processing, plus optional discount programs (e.g., **SageMaker Savings Plans**, introduced in 2021) and a free tier.

Get Vitaly Sem's stories in your inbox
Join Medium for free to get updates from this writer.
Enter your email
Subscribe

Gotchas the author calls out:

- **Ancillary charges** add up: CloudWatch (≈$0.10 per 1K metrics), EBS (Elastic Block Store) volumes for notebooks, inter-AZ (Availability Zone) data transfer (~$0.01/GB). These are easy to miss in early budgets.
- **Savings Plan trade-offs:** commitment is tied to specific instance families; a quoted **$547/year** minimum can be uneconomical for small projects and hard to reallocate if your workload shifts.
- **Free tier limits: ~250 hours** on `ml.t2.medium`—about **10 days** nonstop—and **no GPU** access. Useful to click around the UI, not to run meaningful modern

training.

- **SageMaker Studio** (≈$24/month) unlocks experiment tracking, data UX, and collaboration. Without it, the free tier feels too bare for pro use.

Net: SageMaker shines economically **once workloads are steady and predictable**. In early-stage R&D, hidden line items can push actuals **30–50%** above the napkin math.

Flexibility & developer experience

- **Custom Docker** is first-class: bring nightly PyTorch, custom TensorFlow, etc. Expect longer image builds (reported **up to ~15 minutes** vs. 2–3 for simple bases).
- Distributed frameworks like **Ray** and **Dask** are supported; Dask clusters often need manual setup via Processing jobs.
- Dependencies are pinned and built at image-creation time, which helps avoid “works on my machine” drift.

Speed to production

- Commit-to-pipeline typically **10–15 minutes** due to mandatory Docker build steps even for small code changes.
- **SageMaker Projects** wires CI/CD (Continuous Integration/Continuous Delivery) via CodePipeline/CodeBuild, but you’ll need some AWS fluency.

Strength: elastic endpoints

Autoscaling for inference endpoints is excellent — smoothly handling **hundreds to tens of thousands of requests per minute** without manual tuning.

Ecosystem lock-in (the double-edged sword)

- Top-tier integrations with **S3**, **CloudWatch**, **IAM (Identity and Access Management)**, **Redshift**, **Glue**.
- **SageMaker Pipelines** offer mature MLOps, but introduce a **DSL (Domain-Specific Language)** and deeper coupling to AWS semantics.

Google Cloud Platform (GCP) — Vertex AI

Google's **Vertex AI** favors opinionated, integrated ML development. Compared to AWS's "constructor set," Vertex AI pushes **ready-to-use cycles** and tight data integrations.

Pricing posture

- Optimized serving (e.g., via pre-built containers) and **batching** can bring **prediction cost** down substantially; for classic vision models like **ResNet-50**, figures as low as **~$0.0001/request** are cited in tuned setups.
- **Per-second billing** on **TPUs (Tensor Processing Units)** (e.g., **$1.35/hour** for TPU v2) reduces overpay for short tests.
- **Sustained Use Discounts:** automatic reductions after 25% monthly utilization (e.g., **V100** dropping from ~$2.48/h to ~$1.74/h) without upfront commitments.

Hidden costs to watch:

- **Network egress/ingress:** inter-region up to **$0.19/GB**; public egress ~**$0.08/GB**. Heavy load-balancing can add **15–20%** to bills.
- **Ops telemetry:** Logging ~**$0.50/GB**; tracing/profiling can rack up hundreds per month on large models.
- **Storage & analytics: BigQuery** adds ~$0.02/GB-month plus per-query costs;

ML prediction queries (e.g., $0.10 per 1K) scale quickly.

- **Idle resources:** Stopped but **not fully released** GPUs/VMs can silently accrue charges — an easy $3,000/month mistake.

Operational advice: automate resource lifecycles, enforce budgets/quotas, and audit network flows. Otherwise, **network + observability** can eclipse compute by **1.5–2×**.

ML features that stand out

- **Vertex AI + BigQuery**: train on warehouse data without complex ETL (Extract, Transform, Load).
- **Feature Store** keeps training/serving features in sync; **Vertex Pipelines** gives lineage, experiment tracking, versioning.
- **TPUs**: for large language models and vision/NLP, TPUs can cut training time dramatically (e.g., large BERT under an hour on TPUs vs. >5 hours on 8× V100 GPUs). You can **mix TPU and GPU** per stage.

Microsoft Azure Machine Learning (Azure ML)

Azure ML excels in **enterprise rollouts**: security, governance, and **hybrid** deployment are first-class. Train in cloud, **serve on-prem** (on-premises) if needed — great for data residency and regulatory controls.

Pricing posture

- **Platform fee: ~$0.075/hour** for base ML workspace features; free tier includes **10 GB** model storage and basic monitoring.
- **GPU pricing: NVIDIA V100** around **$2.55/hour** (region-dependent).
- **Distributed CPU (Central Processing Unit):** from **$0.045 per vCPU-**

hour.

- **Spot VMs** (capacity-reclaimable) can cut compute **up to ~90%** with preemption-tolerant jobs.

Hidden costs to watch:

- **Modular billing** across services stacks up:
- **AKS (Azure Kubernetes Service)** control plane (**~$0.10/hour**) for model serving
- **Azure Storage** (**~$0.018/GB-month**)
- **Inter-region data transfer** (**~$0.02/GB**) even on private backbone
- The upside: for organizations on **Microsoft 365/Office 365**, **Active Directory**, and **Power BI**, the integration and SSO (Single Sign-On) story is hard to beat.

Enterprise strengths

- Tight hooks into **Power BI** (analytics), **Azure Data Lake** (big data), **Defender for Cloud** (security).
- **Hybrid/edge** via **Azure IoT Edge**: deploy models to factory/retail/field gear, sync results to cloud.
- End-to-end automation: dataset versioning, lineage, model monitoring in production.

When each platform makes the most sense

Platform Best For Why **AWS SageMaker** Expert teams needing **fine-grained control** and **elastic at-scale inference** Deep AWS integrations, custom Docker

everywhere, robust endpoint autoscaling **GCP Vertex AI Research & startups** optimizing for **time-to-insight** and **warehouse-native ML** BigQuery + Vertex AI, TPUs for rapid training, strong managed MLOps defaults **Azure ML Enterprises** with Microsoft stacks, **compliance** and **hybrid/on-prem** needs AD-driven security, Power BI, IoT Edge, predictable enterprise onboarding

A pragmatic decision framework

1. **Data gravity:** If most data lives in S3 → **SageMaker**; BigQuery → **Vertex AI**; Azure Data Lake / SQL Server → **Azure ML**.
2. **Training profile:**

- Large-model pretraining / LLM/NLP vision → consider **TPUs** (**Vertex AI**).
- Heavy custom containers and exotic stacks → **SageMaker**.
- On-prem inference or strict residency → **Azure ML**.

3. Ops maturity:

- Greenfield team: prefer opinionated pipelines (**Vertex AI**).
- AWS-savvy DevOps: **SageMaker Projects** + CodePipeline.
- Corporate IT + SSO/compliance: **Azure ML**.

4. Cost controls:

- Enforce budgets, auto-shutdown, and tag-based chargeback across all clouds.
- Model serving: benchmark with and without batching; measure real egress and observability volume.

TL;DR

- **There's no universal "best."** Match the platform to your **data location**, **team skills**, **regulatory constraints**, and **latency/throughput** goals.
- **AWS SageMaker** = maximum flexibility and industrial-grade endpoint scaling — great once workloads stabilize.
- **GCP Vertex AI** = fastest path from dataset to results (BigQuery + Vertex, TPUs) with sensible managed defaults.
- **Azure ML** = enterprise integration, security, and hybrid deployments — ideal if you're already a Microsoft shop.
- Whatever you pick, **model the hidden costs** (network, logs/metrics, storage) and set up **guardrails** on day one.

AWS

Gcp

Azure

Cloud Computing

Artificial Intelligence

Written by Vitaly Sem

55 followers

·

22 following

Born2Code

No responses yet

Write a response

What are your thoughts?

Cancel

Respond

More from Vitaly Sem

Enabling Internet Access for LM Studio Models. Practical Guide

Vitaly Sem

Enabling Internet Access for LM Studio Models. Practical Guide

I'm a huge fan of LM Studio because it lets you run powerful AI models locally on your machine—giving you maximum privacy for your…

Aug 4, 2025

A clap icon7

A response icon6

Django 6.0 Background Tasks. First Look

Vitaly Sem

Django 6.0 Background Tasks. First Look

Django 6.0 finally ships with a built-in framework for background tasks: django.tasks.

Dec 9, 2025

A clap icon2

Error-Handling Patterns in Go Every Developer Should Know

Vitaly Sem

Error-Handling Patterns in Go Every Developer Should Know

Error handling in Go (Golang) is more than preventing crashes — it's how you build systems that fail gracefully, explain what went wrong…

Oct 21, 2025

A clap icon 8

A response icon 1

Python Generators: The Power of send, throw, and close. Tutorial

Vitaly Sem

Python Generators: The Power of send, throw, and close. Tutorial

In my previous article Python Generator Functions Explained, we explored how to create generator objects and control their execution using…

Aug 23, 2025

A clap icon 1

See all from Vitaly Sem

Recommended from Medium

In
Stackademic
by
Tushar

Real-World Kubernetes Interview Questions

Production Best Practices & Anti-Patterns

Jan 24

A clap icon63

Zero ETL Is the Reality Check Every Data Engineer Needs in 2026

AWS

In

AWS in Plain English

by

Khushbu Shah

Zero ETL Is the Reality Check Every Data Engineer Needs in 2026

Why the best data engineers are becoming platform designers?

Jan 19

A clap icon174

A response icon4

Building AI Agents in 2026: Chatbots to Agentic Architectures

In

Level Up Coding

by

Gaurav Shrivastav

Building AI Agents in 2026: Chatbots to Agentic Architectures

This is the engineering blueprint for building production-ready agentic systems that actually work.

Jan 22

A clap icon406

A response icon5

I Tried ClawdBot for Data Engineering and Here’s the Honest Truth

Reliable Data Engineering

I Tried ClawdBot for Data Engineering and Here’s the Honest Truth

The AI agent that runs on your server sounds amazing in theory. In practice? It’s complicated.

Jan 26

A clap icon56

A response icon3

Anthropic Just Released Claude Code Course (And I Earned My Certificate)

In

AI Software Engineer

by

Joe Njenga

Anthropic Just Released Claude Code Course (And I Earned My Certificate)

Anthropic just launched their Claude Code in Action course, and I've just passed —how about you?

Jan 21

A clap icon 1.7K

A response icon 20

10 Essential Docker Concepts I Wish Someone Explained to Me on Day One

Cloud With Azeem

10 Essential Docker Concepts I Wish Someone Explained to Me on Day One

When I first started learning Docker, everyone told me the same thing: “It’s simple. Just learn images and containers.” What they didn’t…

6d ago

A clap icon143

See more recommendations